

Stream Processing Systems

Apache Spark and Spark Streaming

Objectives


- Introduction to Streaming Applications
 - Core development concepts
- Apache Spark, SparkSQL and Spark Streaming
 - Provide enough overview to get started
- Programming Hands On
 - <https://spark.apache.org/>

Anatomy of a Stream Processing System

Integrated Development Environment

Scale-Out Runtime

Analytic Toolkits



Development and Management

Flexibility and Scalability

Functional and Optimized

Stream Processing Middleware Services

- Distribution
 - Plumbing related to moving processes around
- Resource management
 - Placement, scheduling, load shedding
- Job management
 - Starting/stopping/editing jobs, handling dynamic connections
- System monitoring
 - Health, state, performance of the system
- Security
- Logging

- Data Transport
 - Inter-Process, intra-Process, intra-host, inter-host
- Debugging
- Visualization


Stream Processing Applications

- Composition Language
- Defining Tuples and Streams
- Stream Sources
- Operators and Stream Transformers
- Windows and Stream Relational Processing
- Stream Sinks

Structure of a Streaming Application

SPL Hello World: A fully Declarative Model

```
1 namespace sample;
2
3 composite HelloWorld {
4   graph
5     stream<uint32 id> Beat = Beacon() {
6       param
7         period: 1.0; // seconds
8         iterations: 10u;
9       output
10        Beat: id = (uint32) (10.0*random());
11    }
12    stream<rstring greeting, uint32 id> Message = Functor(Beat) {
13      output
14        Message: greeting = "Hello World! (" + (rstring) id + ")";
15    }
16    () as Sink = FileSink(Message) {
17      param
18        file: "/dev/stdout";
19        format: txt;
20        flush: 1u;
21    }
22 }
```



Apache Spark and Spark Streaming

Apache Spark History

- Open-source in-memory cluster computing
- Developed at AMPLab
 - UC Berkeley (2009)
 - Open sourced in 2010
- First release: May 2014
- Overcomes several limitations of Hadoop Map-Reduce model

Apache Spark Platform

**Development
Environment**

Scale-Out Runtime

Analytic Toolkits



Python, Scala, Java

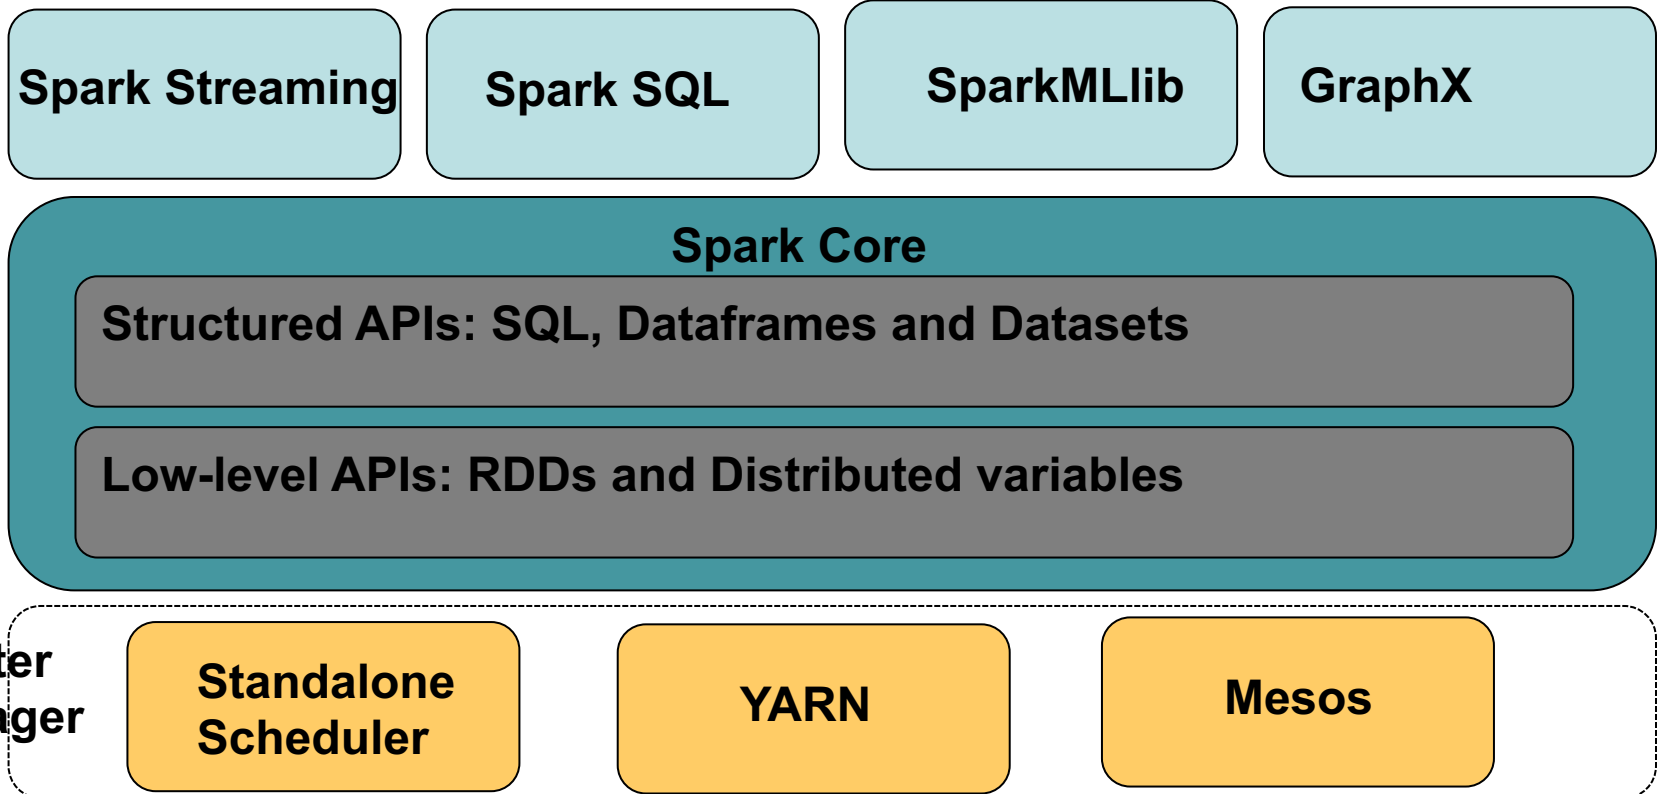


In-Memory Cluster Computing

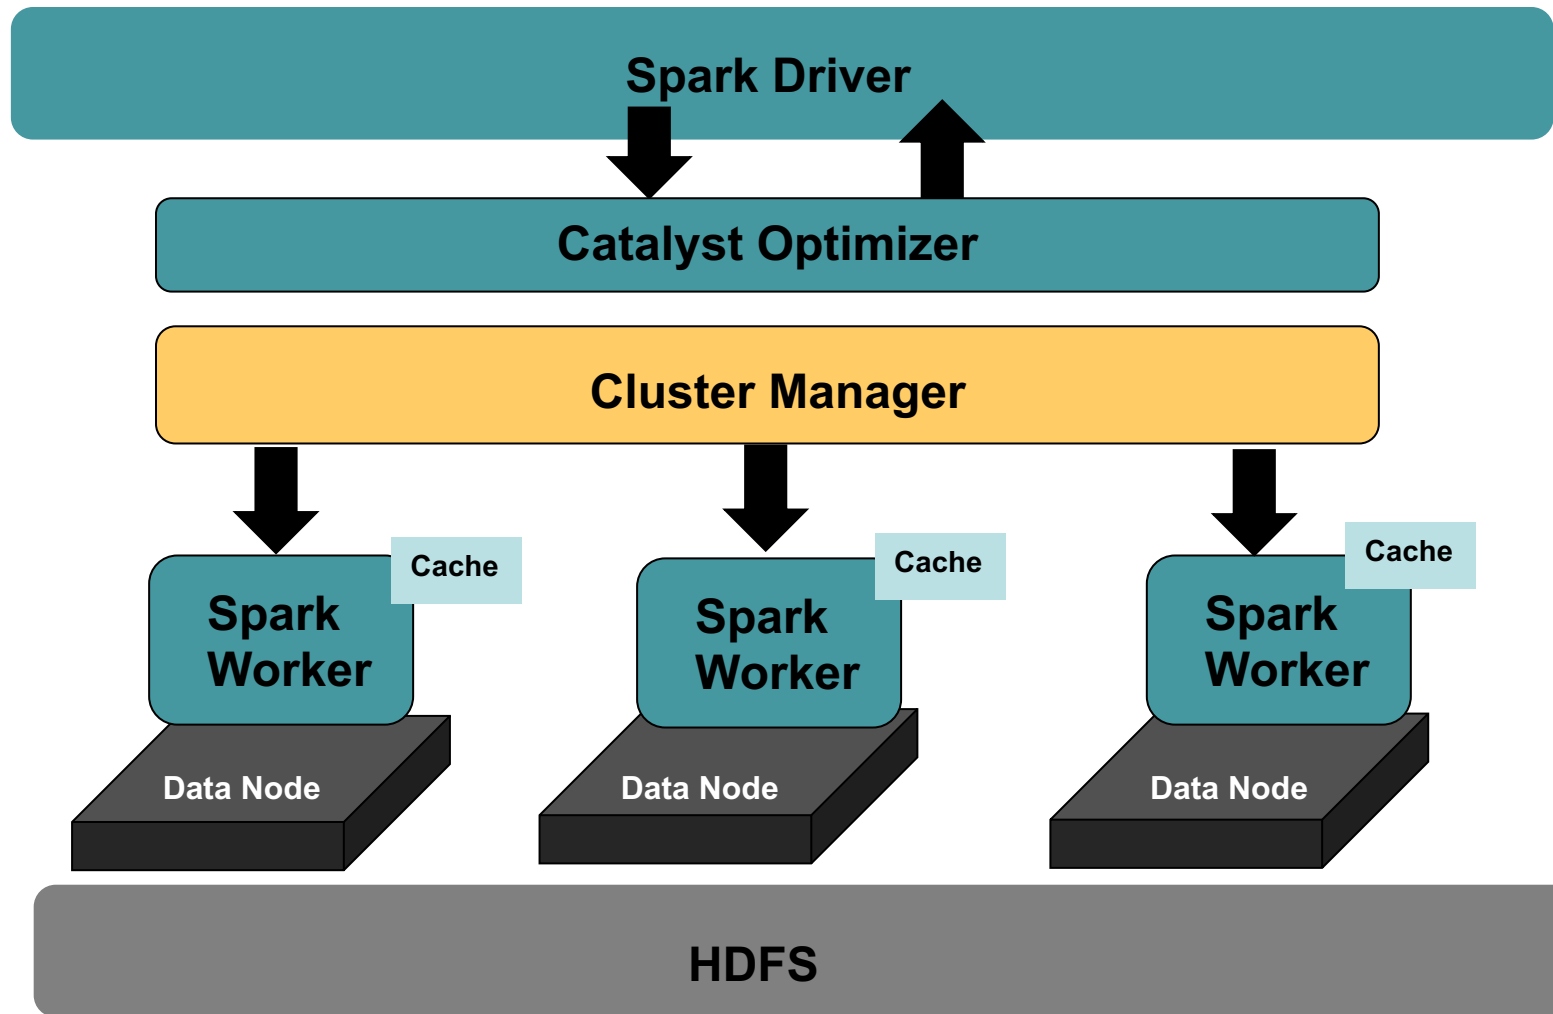


SparkMLlib, Connectors

Spark Components



Spark Runtime Architecture



Recall: Middleware Support

- Access transparency
 - Local and remote components accessible via same operations
- Location transparency
 - Components locatable via name, independent of location
- Concurrency transparency
 - Components/objects can be run in parallel
- Migration transparency
 - Allows movement of components without affecting other components
- Replication transparency
 - Allows multiple instances of objects for improved reliability (mirrored web pages)
- Failure transparency
 - Components designed while accounting for failure of other services

Spark Basics

- **Installation**

- Straightforward on Linux/Windows
- Download binaries – unzip
 - <http://spark.apache.org/downloads.html>
- `pip install pyspark`

- **Includes**

- `Readme.md`
- `bin` (binaries, including shell)
- `core, streaming, python` (source code)
- `examples`

Spark Shell (Python and Scala)

- Shell
 - `pyspark` (Python)
 - `spark-shell` (Scala)
- Interactive command line based interface
- Similar to those in R, Python, Matlab
- Log level can be controlled
 - `log4j.properties.template` in `conf` folder
 - Copy to `log4j.properties` and set log level
 - `log4j.rootCategory=WARN`
- Easy to get started

Resilient Distributed Datasets (RDDs)

- Resilient
 - Data can be recreated if lost
- Distributed
 - Stored in memory across the cluster
- Dataset
 - Immutable
 - Lazy Evaluated
 - Cacheable
 - Type Inferred
- Fundamental unit of processing in Spark
 - Dstreams for Streaming

Immutability

- Data cannot be changed
 - Allows improved parallelism and scalability
 - However requires more storage

```
const int a = 0; //immutable
int b = 0; //mutable
b++; //allowed in-place change
a++; //not allowed
c = a+1; //create new variable - leave old variable intact
```

- Every transformation creates a new copy!

Lazy Evaluation

- Transformations are not computed till needed
 - Deferred evaluation
 - Separate execution from evaluation

```
C2 = C1.map(lambda val: val+1) //not evaluated
C3 = C2.map(lambda val: val+2) //not evaluated
print C3 //now evaluated as
```

```
C3 = C1.map(lambda val : ((val+1)+2))
```

- Multiple transformations performed together when result needs to be reported
- Laziness works only when immutable data

Type Inference

- Type of a variable/collection inferred from values by compiler
- Sometimes can be inferred from transformation applied
 - Some transformations have fixed return type

```
C1 = [1, 2, 3, 4, 5] //Type array (explicit)
C2 = C1.map(lambda val: val+1) //C2 type array
C3 = C2.count() //Type integer
```

Caching

- Immutable data can be cached over long periods
- Lazy transformations allow for data to be recreated as needed
- Transformations can also be saved
- Caching improves performance

Programming with Spark

- Hands on Hello World

```
//read file line by line to create RDD called lines
lines = sc.textFile("filename")
//return count of lines in lines RDD
lines.count()
//return the first element of lines - first line in file
lines.first()
```

- SC: Spark Context

- Driver programs access Spark runtime through this context object
- Represents a connection to the Spark Cluster
- Created automatically in the shell

Programming with Spark

- Hands on Hello World

```
//read file line by line to create RDD called lines
lines = sc.textFile("filename")
//filter the RDD to extract lines with text test in them
filteredLines = lines.filter(lambda val: "test" in val)
//return count of lines in lines RDD
lines.count()
//return the first element of filteredLines
filteredLines.first()
```

- filter
 - Special type of transformation to create a new RDD
 - Recall immutability and lazy transformations

Programming with Spark

- **lambda Syntax**

- Allows for definition of inline functions in Python on Spark

```
//filter the RDD to extract lines with text test in them
filteredLines = lines.filter(lambda val: "test" in val)
```

- Defines a boolean text filter in this case – that is applied to all the data items within the RDD. Can also be defined explicitly

```
//filter the RDD to extract lines with text test in them
def findTest(x):
    return "test" in x
```

```
filteredLines = lines.filter(findTest)
```

- Spark ships this function to all the executor nodes so that it can be run in parallel
- Equivalent => syntax for Scala

Standalone Programs with Spark

- Can invoke a program from outside the shell
 - Allows for ease of programming

```
//use spark-submit to submit custom python code  
spark-submit mycode.py
```

```
from pyspark import SparkConf, SparkContext  
sc = SparkContext("local", "myApp")  
lines = sc.textFile("foo")  
filteredLines = lines.filter(lambda x: "test" in x)
```

- In mycode.py “local” Master → single node deploy
 - Use a cluster URL that tells the driver how to connect to the cluster

Programming with Spark

- Two types of operations on RDDs
 - Transformations, Actions
- Transformations
 - Create new RDD by applying transformation to the RDD entries
- Actions
 - Compute a result based on an RDD and return to the driver program or write to output, e.g. HDFS
- Important distinction
 - Lazy evaluation model of Spark
 - Need for `persist` of RDDs: RDDs are recomputed every time an action is run on them. `persist` allows for caching of the result so that intermediate results can be reused

Programming with Spark

- RDD Creation

```
//parallelize method  
lines = sc.parallelize(["My", "test", "example here"])  
//read from file  
lines2 = sc.textFile("filename")
```

- RDD Transformations

```
//filter  
filteredLines = lines.filter(lambda x: "test" in x)  
filteredLines2 = lines.filter(lambda x: "example" in x)  
  
//union  
final = filteredLines.union(filteredLines2)
```

- What happens if we do not persist? What about lazy evaluation? ORs and ANDs?

Programming with Spark

- Element wise transformations
 - `filter` and `map`
- `map`
 - Apply stateless transform to each element within RDD

```
lines = sc.parallelize([1, 4, 3])  
//read from file  
lines2 = lines.map(lambda x: x*x)  
res = lines2.collect()  
for num in res:  
    print num
```

- Element wise squaring of all the input values
- What does `collect` do?
 - Brings all the results back to the driver program as an array
 - Important to know – when RDD distributed across multiple nodes

Other Element Wise Transformations

- Element-wise one to many transformation
 - flatMap
 - Can create more (or less) than one element in output RDD per element in input RDD

```
//parallelize method
lines = sc.parallelize(["My", "test", "example here"])
//parse the text into words using space separator
words = lines.flatMap(lambda x: x.split(" "))
```

- Third element in `lines` gets mapped to more than one element in the `words` RDD
- What is the difference from?

```
//parallelize method
tokenLines = lines.map(lambda x: x.split(" "))
```

Set Operations on RDDs

- `union`
 - Set-wise union of two RDDs – retains duplicates
- `distinct`
 - Remove duplicate values in RDD
 - Expensive operation as needs to aggregate across all partitions of the RDD
- `subtract`
 - Set-wise difference between RDDs
 - Also removes duplicates → expensive operation
- `intersection`
 - Set-wise common elements between RDDs
 - Also removes duplicates

Set Operations on RDDs

```
lines = sc.parallelize(["coffee", "coffee", "panda", "monkey",  
"tea"])  
lines2 = sc.parallelize(["coffee", "monkey", "kitty"])  
  
res1 = lines.union(lines2)  
//{"coffee", "coffee", "coffee", "panda", "monkey", "monkey", "tea",  
"kitty"}  
  
res2 = lines.distinct()  
//{"coffee", "panda", "monkey", "tea"}  
  
res3 = lines.intersection(lines2)  
//{"coffee", "monkey"}  
  
res4 = lines.subtract(lines2)  
//{"panda", "tea"}
```

Set Operations on RDDs

- cartesian operation
 - Create cartesian product of sets

```
lines = sc.parallelize(["coffee", "coffee", "panda"])
lines2 = sc.parallelize(["coffee", "monkey"])
```

```
res1 = lines.cartesian(lines2)
//{("coffee", "coffee"), ("coffee", "monkey"), ("coffee", "coffee"),
  ("coffee", "monkey"), ("panda", "coffee"), ("panda", "monkey")}
```

- Can be very expensive for large sets

Summary of Basic RDD Operations

```
lines = sc.parallelize([1, 2, 3, 3])
```

Name	Purpose	Example	Result
<code>map()</code>	Element-wise function, same number of output elements	<code>lines.map(lambda x: x*x)</code>	{1, 4, 9, 9}
<code>flatMap()</code>	Element-wise function, different number of output elements	<code>lines.flatMap(lambda x: [0,x])</code>	{0,1,0,2,0,3,0,3}
<code>filter()</code>	Select RDD elements that satisfy a condition	<code>lines.filter(lambda x: x>1)</code>	{2, 3, 3}
<code>distinct()</code>	Remove duplicates - expensive	<code>lines.distinct()</code>	{1, 2, 3}
<code>sample(withReplacement, fraction, [seed])</code>	Randomly sample from RDD	<code>lines.sample(false, 0.5)</code>	{1,3} Non-deterministic

Summary of Two-RDD Operations

```
a = sc.parallelize([1, 2, 3])
```

```
b = sc.parallelize([3, 4, 5])
```

Name	Purpose	Example	Result
<code>union()</code>	Union across all elements – duplicates retained	<code>a.union(b)</code>	{1, 2, 3, 3, 4, 5}
<code>subtract()</code>	Set-difference – duplicates removed	<code>a.subtract(b)</code>	{1, 2}
<code>intersection()</code>	Common elements – duplicates removed	<code>a.intersection(b)</code>	{3}
<code>cartesian()</code>	Cartesian product - expensive	<code>a.cartesian(b)</code>	{(1,3), (1,4), (1,5), (2, 3), (2,4), (2,5), (3,3), 3,4), (3,5)}

Actions on RDDs: Aggregations

- Produce result of running an operation on RDD
 - not new RDD
 - Often recursive application of some function
- `reduce` action
 - Takes function that operates on two elements on RDD and returns element of same type
 - Applies function recursively across all elements in RDD

```
nums = sc.parallelize([1, 2, 3, 4])  
res = nums.reduce(lambda x, y: x+y)  
//returns sum of entire RDD = 10
```

- Note that `res` is not a RDD – it is an integer value

Actions on RDDs: Aggregations

- Similar to `reduce` is the `fold` action
 - Except we can add an initial value for the result

```
nums = sc.parallelize([1, 2, 3, 4])
res = nums.reduce(lambda x, y: x+y)
//returns sum of entire RDD = 10
res2 = nums.fold((5), (lambda x,y:x+y))
//returns the sum starting with an initial value 5 = 15
```

- More complex action is `aggregate`
 - Frees the constraint on the function that return value has to be same as the element value
 - Takes initial value, function to combine elements into accumulator, and another function to combine accumulators

Actions on RDDs: Aggregations

- Example to compute sum and count at same time
 - Initialize with 0,0

```
nums = sc.parallelize([1, 2, 3, 4])
res = nums.aggregate(
    (0,0),
    (lambda acc, val: (acc[0]+val, acc[1]+1)),
    (lambda acc1, acc2: (acc1[0]+ acc2[0], acc1[1] + acc2[1])))
//returns sum and count = (10,4)
//can compute average
print res[0]/res[1]
```

- First function sums RDD values and increments counter by 1
 - Next function combines the accumulators – potentially across multiple partitions of the RDD
- Can we do something equivalent with map and reduce functions only?

Actions on RDDs: Collection

- The `collect` action
 - Collects all values across the RDD across all partitions and workers and brings them back as a list to driver program (can be expensive for a large RDD)
 - Expects the resulting list to fit in memory
- The `take(n)` action
 - Takes `n` values from the RDD and returns to the driver program as a list
 - Order of return values may vary based on optimization across partitions
- The `top` action
 - Return top value based on a specified order
- The `takeSample(withReplacement, num, seed)` action
 - Return `num` samples using random sampling – with or without replacement

Actions on RDDs

```
lines = sc.parallelize([1, 2, 3, 3])
```

Name	Purpose	Example	Result
<code>collect()</code>	Collect all the RDD elements into a list	<code>lines.collect()</code>	<code>{1, 2, 3, 3}</code>
<code>count()</code>	Count elements in RDD	<code>lines.count()</code>	4
<code>countByValue()</code>	Count based on value of RDD – expensive action	<code>lines.countByValue</code>	<code>{{(1,1),(2,1),(3, 2)}</code>
<code>take(num)</code>	Collect num elements into a list	<code>lines.take(2)</code>	<code>{1, 2}</code>
<code>top(num)</code>	Collect num elements based on order	<code>lines.top(2)</code>	<code>{3,3}</code>

Actions on RDDs

```
lines = sc.parallelize([1, 2, 3, 3])
```

Name	Purpose	Example	Result
<code>takeOrdered(num, (ordering))</code>	Collect num elements based on order function	<code>lines.takeOrdered(2, (myorderingfunc))</code>	{1, 2}
<code>takeSample(withReplacement, num, [seed])</code>	Get elements in RDD	<code>lines.takeSample(false, 2)</code>	Non-deterministic
<code>reduce(func)</code>	Recursive aggregation of RDD based on function	<code>lines.reduce(lambda x, y: x+y)</code>	9
<code>fold(init, func)</code>	Recursive aggregation of RDD based on function with init val	<code>lines.fold(2, (lambda x, y: x+y))</code>	11
<code>Aggregate(init, seqOp, combOp)</code>	Recursive aggregation that allows combination function with return type different than RDD element type	See earlier example	
<code>foreach(func)</code>	Apply function for each element of RDD	<code>lines.foreach(func)</code>	Depends on func

RDD Persistence

- Due to lazy evaluation
 - RDDs are processed only when action applied
 - If multiple actions applied to same RDD – it may be computed multiple times
 - Persistence allows for caching of intermediate results

```
nums = sc.textFile("foo")
res = nums.top()
res2 = nums.count()
//File read two times
```

```
nums = sc.textFile("foo")
nums.persist()
res = nums.top()
res2 = nums.count()
//File read only once
```

- Each node stores its computed partition for reuse
- If a node fails, the partition is recomputed
- Old partitions evicted using a Least Recently Used (LRU) policy

Pair RDDs

- RDDs that have key,value elements
 - Common to a majority of data analysis tasks
 - Expose new operations and partitioning
 - Can be read from certain data formats
 - Created using the map function

```
lines = sc.textFile("foo")  
pairlines = lines.map(lambda x: (x.split(" ")[0], x))
```

- Taking the first word in a line as the key to the line
- This returns tuple elements within the resulting RDD
- Have additional transformation functions that can be applied

Pair RDD Transformations

Rdd: `{(1,2), (3,4), (3, 6)}` lines: `{("a","This is"), ("b", "a test")}`

Name	Purpose	Example	Result
<code>groupByKey()</code>	Group elements in RDD by key to make element lists in new RDD	<code>Rdd.groupByKey()</code>	<code>{{(1,[2]),(3,[4,6])}}</code>
<code>mapValues(func)</code>	Apply function to each value without changing the key	<code>Rdd.mapValues(lambda x: x*x)</code>	<code>{{(1,4),(3,16),(3,36)}}</code>
<code>flatMapValues(func)</code>	Apply function to each RDD element without changing the key – each function can return more than one value	<code>lines.flatMapValues(lambda x:x.split(" "))</code>	<code>{{("a","This"), ("a","is"),("b","a"),("b","test")}}</code>
<code>keys()</code>	Return an RDD with just the keys	<code>Rdd.keys()</code>	<code>{1,3,3}</code>

Pair RDD Transformations

Rdd: $\{(1, 2), (3, 4), (3, 6)\}$ other: $\{(3, 9)\}$

Name	Purpose	Example	Result
<code>values()</code>	Return an RDD with just the values	<code>Rdd.values()</code>	$\{2, 4, 6\}$
<code>sortByKey()</code>	Create new RDD sorted by key	<code>Rdd.sortByKey()</code>	$\{(1, 2), (3, 4), (3, 6)\}$
<code>subtractByKey()</code>	Set operation - remove elements with key common to second RDD . Removes duplicates	<code>Rdd.subtractByKey(other)</code>	$\{(1, 2)\}$
<code>cogroup()</code>	Group elements across two RDDs with the same key. Also does a <code>groupByKey</code> within each RDD	<code>Rdd.cogroup(other)</code>	$\{(1, [2, []]), (3, [4, 6], [9])\}$

Join: A Database Perspective

- An SQL join on tables in a relational database :

```
SELECT NAME, EQUIP
```


```
FROM EMP-DEPT JOIN DEPT-EQUIP ON
```

```
EMP-DEPT.DEPT = DEPT-EQUIP.DEPT
```

- Inner Join*: Result contains only tuples that were matched

EMP-DEPT

Name	Dept.
Alice	Accounting
Bob	Programming
Carol	Sales
Dan	HR
Ellen	Programming
Frank	Accounting
George	Research



DEPT-EQP

Dept.	Equipment
Accounting	Calculator
Programming	Computer
Sales	Computer
HR	Tablet
Security	Vehicle
Sales	Phone



Name	Equipment
Alice	Calculator
Bob	Computer
Ellen	Computer
Dan	Tablet
Carol	Computer
Carol	Phone
Frank	Calculator

Join: A Database Perspective

- An SQL join on tables in a relational database :

```
SELECT NAME, EQUIP
```


```
FROM EMP-DEPT LEFT OUTER JOIN DEPT-EQUIP ON
```

```
EMP-DEPT.DEPT = DEPT-EQUIP.DEPT
```

- Left Outer* Join: Result contains tuples that were matched, and unmatched tuples from the “left” table

EMP-DEPT

Name	Dept.
Alice	Accounting
Bob	Programming
Carol	Sales
Dan	HR
Ellen	Programming
Frank	Accounting
George	Research



DEPT-EQP

Dept.	Equipment
Accounting	Calculator
Programming	Computer
Sales	Computer
HR	Tablet
Security	Vehicle
Sales	Phone



Name	Equipment
Alice	Calculator
Bob	Computer
Ellen	Computer
Dan	Tablet
Carol	Computer
Carol	Phone
Frank	Calculator
George	NULL

Join: A Database Perspective

- An SQL join on tables in a relational database :

```
SELECT NAME, EQUIP
```


```
FROM EMP-DEPT RIGHT OUTER JOIN DEPT-EQUIP ON
```

```
EMP-DEPT.DEPT = DEPT-EQUIP.DEPT
```

- Right Outer Join*: Result contains tuples that were matched, and unmatched tuples from the “right” table

EMP-DEPT

Name	Dept.
Alice	Accounting
Bob	Programming
Carol	Sales
Dan	HR
Ellen	Programming
Frank	Accounting
George	Research



DEPT-EQP

Dept.	Equipment
Accounting	Calculator
Programming	Computer
Sales	Computer
HR	Tablet
Security	Vehicle
Sales	Phone




Name	Equipment
Alice	Calculator
Bob	Computer
Ellen	Computer
Dan	Tablet
Carol	Computer
Carol	Phone
Frank	Calculator
NULL	Vehicle

Aside: Additional Joins in Spark

- Semi-joins: Left and Right
Only return rows from the left (or right) table that match, and no duplicates

EMP-DEPT

Name	Dept.
Alice	Accounting
Bob	Programming
Carol	Sales
Dan	HR
Ellen	Programming
Frank	Accounting
George	Research



DEPT-EQP

Dept.	Equipment
Accounting	Calculator
Programming	Computer
Sales	Computer
HR	Tablet
Security	Vehicle
Sales	Phone



Name	Equipment
Alice	Accounting
Bob	Programming
Ellen	Programming
Dan	HR
Carol	Sales
Frank	Accounting

- Anti-joins: Apply a not to a Semi-join
 - Left and Right

Joins on Pair RDDs

- Joins operate like database counterparts
 - Inner, LeftOuter and RightOuter
 - Equality conditions are defined based on keys

```
Rdd: { (1,2), (3,4), (3, 6) }  other: { (3,9) }
```

```
Rdd.join(other): { (3, (4,9)), (3, (6,9)) }
```

```
Rdd.leftOuterJoin(other): { ((1, (2,None)), (3, (4, 9))),  
                           (3, (6, Some(9))) }
```

```
Rdd.rightOuterJoin(other): { (3, (4,9)), (3, (6,9)) }
```

- Is Join a symmetric operation?
- How about RightOuterJoin and LeftOuterJoin?
- How about non-equi joins?

Aggregations on Pair RDDs

- **reduce, fold and aggregate actions extended to work by key**
 - `reduceByKey`, `foldByKey`, `aggregateByKey`
 - These operate as transformations rather than as actions, i.e. return RDDs, since each RDD may have multiple keys

```
Rdd: {(1,2), (3,4), (3, 6)} other: {(3,9)}
```

```
Res = Rdd.reduceByKey(lambda x,y:x+y)  
//Res = {(1,2), (3,10)}
```

```
Res2 = Rdd.foldByKey(3,lambda x,y:x+y)  
//Res2 = {(1,5), (3,13)}
```

```
Res3 = Rdd.mapValues(lambda x: (x,1)).reduceByKey(lambda x,y:  
(x[0]+y[0],x[1]+y[1]))
```

- What does the above do?

WordCount Example on Spark

- WordCount
 - Count occurrences of each word within a file

```
lines = sc.textFile("foo")
words = lines.flatMap(lambda x: x.split(" "))
wordMap = words.map(lambda x: (x,1))
wordCount = wordMap.reduceByKey(lambda x,y: x+y)
```

```
wordCount2 = words.countByKey()
```

```
wordCount3 = words.count()
```

- What is the difference in these different results?

Partitioning and Parallelism

- How does Spark parallelize the processing?
 - Every RDD has a fixed number of partitions
 - When grouping or aggregating RDDs, we can specify degree of partitions

- Especially true and useful for pair RDDs

```
Data = sc.parallelize([("a",1), ("b",2), ("c",3)])  
Res = Data.reduceByKey(lambda x,y: x+y)  
//default parallelism  
Res = Data.reduceByKey(lambda x,y: x+y, 10)  
//Custom parallelism use 10 partitions
```

- Sometimes data may need to be repartitioned to improve efficiency of the operation – for groupings

Partitioning and Parallelism

- Optimize partitioning to minimize communication across nodes
 - Especially useful when a dataset is used multiple times
- Spark does not allow explicit control on which worker node a key goes to
 - Can specify which sets of keys go together
 - E.g. hash partition by key mod 100 → all keys with the same remainder after hashing to 100 will be on same node
 - Range partition based on key values – so that keys in same range appear on same node

```
Data = sc.parallelize([("a",1), ("b",2), ("c",3)])  
Res2 = Data.partitionBy(new HashPartitioner(100)).persist()
```

- Scala example
- Custom partitioners can be written in Scala and Java

Other Actions on Pair RDDs

Rdd: `{(1,2), (3,4), (3, 6)}` other: `{(3,9)}`

Name	Purpose	Example	Result
<code>collectAsMap()</code>	Collect all results as a map and return to driver program	<code>Rdd.collectAsMap()</code>	Map object
<code>countByKey()</code>	Count number of elements for each key	<code>Rdd.countByKey()</code>	<code>{{(1,1),(3,2)}}</code>
<code>lookup()</code>	Return values that have a matching key as an array	<code>Rdd.lookup(3)</code>	<code>[4,6]</code>

Next

- Spark Structured APIs and Dataframes

Structured Spark: Dataframes

- Part of the structured API to handle and process data
 - Equivalent to a table with named and typed columns (i.e. with a schema)
 - Just that the table is distributed over multiple workers
- Analogous to Pandas Dataframes

```
linesdf =  
spark.read.option("inferSchema", "true").option("header", "true")  
.csv("data/structured_data.csv")  
  
linesdf.first()
```

- Same notions of lazy execution and actions (like for RDD) carry over

Transformations/Actions on Dataframes

- What do you think this does?

```
from pyspark.sql.functions import desc
```

```
linesdf.groupBy("Line").count().withColumnRenamed("count", "n_products").sort(desc("n_products")).show(5)
```

Transformations on Dataframes

```
df1 = src, dest, flights
```

```
NY    , CA , 100
```

```
CA    , CA , 25
```

```
CA    , NY , 100
```

Name	Purpose	Example	Result
<code>select()</code>	Select columns in dataframe	<code>df1.select("src")</code>	{NY, CA, CA}
<code>selectExpr()</code>	Select and apply transforms and aggregations (can create new cols or dfs)	<code>df1.selectExpr("*", "(src = dest) as inState")</code>	NY, CA, 100, False CA, CA, 25, True CA, NY, 100, False
<code>withColumn()</code>	Explicitly create new columns – can use literals (as constants)	<code>df1.withColumn("const", lit(50))</code>	src, dest, flights, const NY, CA, 100, 50 CA, CA, 25, 50 CA, NY, 100, 50

Transformations on Dataframes

```
df1 = src, dest, flights
```

```
NY    , CA , 100
```

```
CA    , CA , 25
```

```
CA    , NY , 100
```

Name	Purpose	Example	Result
drop()	Drop columns	df1.drop("src")	dest, flights CA, 25 NY, 100 CA, 100
where()	Filter rows	df1.where("flights<50")	CA, CA, 25
distinct()	Unique rows	df1.select ("src").distinct()	NY CA
sort()/order By()	Sort based on some columns	df1.orderBy("src","flights")	src, dest, flights CA, CA, 25 CA, NY, 100 NY, CA, 100
union()	Append dataframe rows with the same schema	df1.union(df1)	

Transformations on Dataframes

```
df1 = src, dest, flights
NY   , CA , 100
CA   , CA , 25
CA   , NY , 100
```

```
def manyFlights(x):
    return x<50
```

Name	Purpose	Example	Result
udf()	Register a user defined function	manyFlightsUdf = udf(manyFlights)	
	Can use this UDF with select or selectExpr	df1.select (manyFlightsUdf(col("flights")))	manyFlights(flights) False True False

Two Dataframe Transforms

```
df1 = src, dest, flights
```

```
NY    , CA , 100
```

```
CA    , CA , 25
```

```
CA    , NY , 100
```

```
df2 = state, pop
```

```
NY    , 100
```

```
CA    , 200
```

Name	Purpose	Example	Result
<code>join()</code>	Applies join across dataframes	<pre>je = (df1.col("src") == df2.col("state")) df1.join(df2, je)</pre>	<pre>src, dest, flights, state, pop CA, CA, 25, CA, 200 CA, NY, 100, CA, 200 NY, CA, 100, NY, 100</pre>
	Support different types of joins		

Actions on Dataframes



Name	Purpose	Example	Result
<code>count()</code>	Count the number of occurrences	<code>df.select(count("src"))</code>	3
<code>countDistinct()</code>	Count number of distinct occurrences	<code>df.select(countDistinct("src"))</code>	2
<code>avg(), sum(), min(), max()</code>	Apply aggregates to column	<code>df.select(max("flights"))</code>	100
<code>approx_count_distinct()</code>	Get approximate distinct counts		
<code>show()</code>	Show a few rows of the dataframe	<code>df1.show()</code>	
<code>collect()</code>	Collect and bring to local memory – show metadata and content	<code>df1.collect()</code>	
<code>take(n)</code>	Take n rows from dataframe – show metadata and content, but flatten	<code>df1.take(5)</code>	

Next

- Spark Streaming

Spark Streaming

- Introduce concept `DStream` (discretized streams)
 - Sequence of RDDs arriving over time
 - Dstreams can be created from different data sources
 - Network interfaces, Flame, Kafka, Flume, File System etc.
 - Support transformations and actions (output operations)




Streaming Example

- Example of constructing Dstream from network socket
 - Setup a Dstream

```
sc = SparkContext(appName="StreamingNetworkWordCount")
ssc = StreamingContext(sc, 1)
lines = ssc.socketTextStream(IP, Port)
counts = lines.flatMap(lambda line: line.split(" "))\
               .map(lambda word: (word, 1))\
               .reduceByKey(lambda a, b: a+b)
counts.pprint()
ssc.start()
ssc.awaitTermination
```

- Take all data received in one second
- Break it into words and then do a wordcount
- Print all results to screen

Streaming Example



The operation of this word count is “stateless” – each window handled independently

Stateless Streaming Transformations

Name	Purpose	Example
<code>map()</code>	Apply function defined inline to each tuple (RDD) in the stream	<code>ds.map(lambda x: x+1)</code>
<code>flatMap()</code>	Apply inline function to each tuple in stream and create tuple with more than one element per input element	<code>ds.flatMap(lambda x: x.split(" "))</code>
<code>filter()</code>	Filter elements in input tuple to create new tuple	<code>ds.filter(lambda x: x.contains("error"))</code>
<code>reduceByKey()</code>	Perform reduce by key within tuple to create new tuple	<code>ds.reduceByKey(lambda (x,y): x+y)</code>
<code>groupByKey()</code>	Group values by key within each tuple	<code>ds.groupByKey()</code>
<code>transform()</code>	Apply any RDD to RDD function on each tuple	<code>ds.transform(lambda rdd: myFunc(rdd))</code>

Stateless Two-Stream Transformations

- Stateless Join on two Streams
 - Perform a standard RDD join on two tuples, one from each stream
 - Fast stream needs to wait for slow stream

```
//Create a streaming context with a 1 second batch
sc = SparkContext(appName="StreamingStatelessJoin")
ssc = StreamingContext(sc, 1)
ipDstream = accessLogs.map(lambda x: (x.getIpAddress(),1))
ipCountDstream = ipDstream.reduceByKey(lambda (x,y): x+y)

ipByteDstream = accessLogs.map(lambda x: (x.getIpAddress(),x.getContentSize()))
ipByteCountDstream = ipByteDstream.reduceByKey(lambda (x,y): (x+y))

ipBCDstream = ipByteCountDstream.join(ipCountDstream)
ssc.start()
ssc.awaitTermination()
```

This can cause memory issues if one stream is much slower than the other
Also – limiting in terms of match conditions, given data arrival order etc,

Stateful Transformation on Streams

- Create and maintain state as tuples processed
- Such state, along with internal algorithm, affects the results
- e.g. DeDuplicate
 - *tuple is considered a duplicate if it shares the same key with a previously seen tuple within a pre-defined period of time*
- Runtime support is challenging
 - Require synchronization in a multi-threaded context
 - No trivial way to parallelize
 - Require some persistence mechanism for fault-tolerance

Windowing

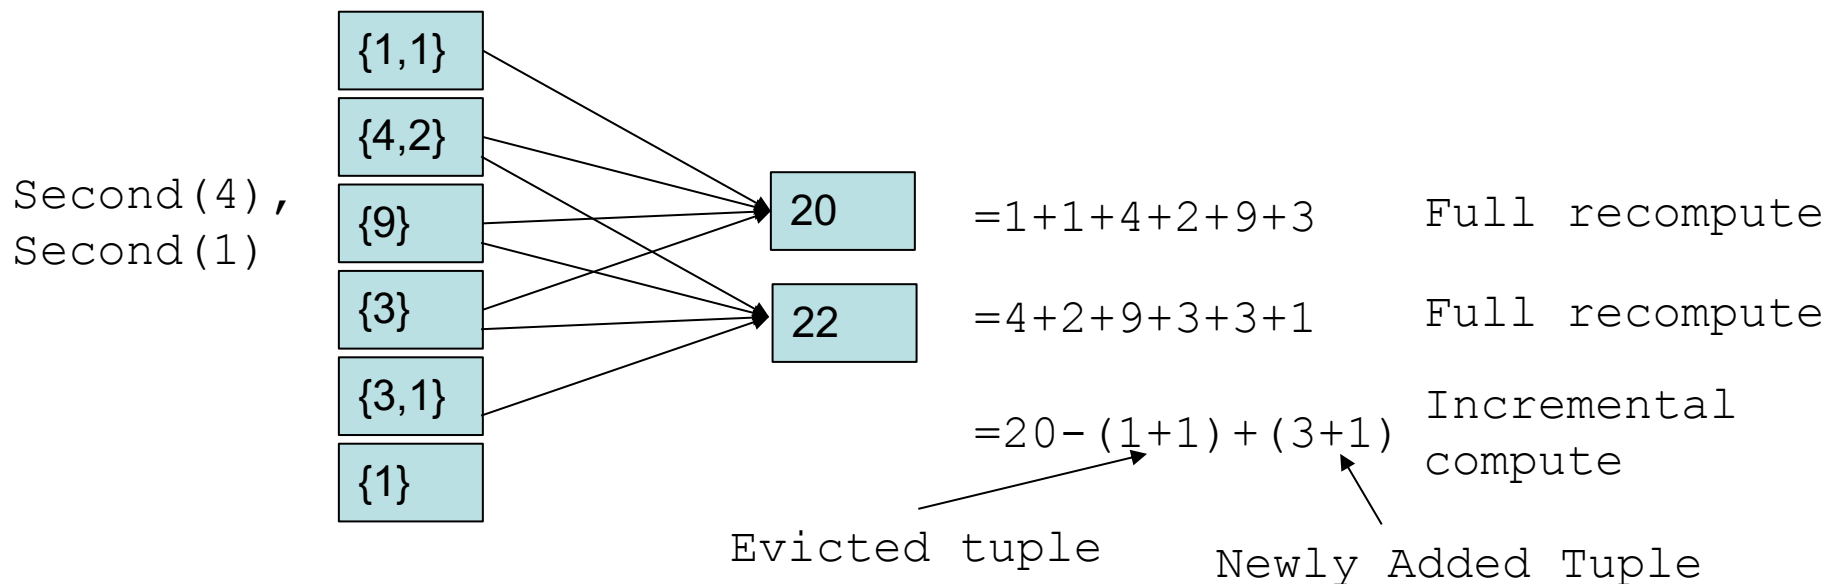
- Window creates a collection of tuples/RDDs within window
- Has size, slide – corresponds to eviction, trigger policies

```
//Create a streaming window  
winLogs = accessLogs.window(30, 10)  
winCounts = winLogs.countByKey()
```

- Slide cannot be smaller than batch size of Dstream
- Performs a Union across RDDs in the window
- Can then apply a reduce across this to compute additional results
 - More efficient aggregations available

Windowed Aggregations

- Incremental operations for aggregations
 - `reduceByWindow`, `reduceByKeyAndWindow`
 - Instead of re-computing fully, remove effect of evicted tuples and add in effect of newly added tuples
 - Require easily reversible functions, e.g. sum



Windowed Aggregations

```
//Create a streaming reduce with a 1 second batch, 4 second window
ipDstream = accessLogs.map(lambda x: (x.getIpAddress(),1))
ipCountDstream = ipDstream.reduceByKeyAndWindow(
    lambda (x,y): x+y,
    4,
    1)
ipCountEDstream = ipDstream.reduceByKeyAndWindow(
    lambda (x,y): x+y,
    lambda (x,y): x-y,
    4,
    1)
```

Reversible function specification followed by window specification allows incremental computation, as opposed to full recompute

Other functions include `reduceByWindow()`, `countByWindow()` and `countByKeyAndWindow()`

Spark Streaming: Inputs and Outputs

- `print()` – grabs first 10 elements from each tuple and prints them to screen
- `saveAsTextFiles("output", "txt")` – writes multiple files into a directory, one per tuple in txt format
- Stream of files
 - `val logData = ssc.textFileStream("logDirectory")`
 - Each file read as a tuple in the Dstream
- Other sources
 - Twitter
 - Kafka, Kinesis, Flume, and others

Limitations of Spark Streaming Model

- All windows defined in wall-clock time
 - Sizes and slides limited to multiples of batch size
- No data-time windows
 - Skewed by randomness in data arrival
 - Cannot handle out of order data
- Limited join and aggregation conditions

Next

- Structured Spark Streaming

References

- H. Karau, A. Konwinski, P. Wendell and M. Zaharia, “Learning Spark”, O’Reilly Press
- <http://spark.apache.org/>
- Chapter 4 in Stream Processing Book